

K N ANANTHA NANDANAN

ananthanandan@outlook.com | linkedin.com/in/de-blank | github.com/ananthanandan | +91 8921397828

TECHNICAL SKILLS

Languages: Python, TypeScript, SQL, BASH, AWK, HTML/CSS

Backend/AI: Django (DRF), FastAPI, Celery, Langchain, LiteLLM, ChromaDB, Pinecone

Infra/LLMops: Docker, RabbitMQ, Redis, Nginx, Coolify, MinIO, Langfuse, Promptfoo, MLflow

Tools/Ops: uv, Git, PostHog, PyMuPDF, OpenCV, Piston API, Linux (Ubuntu/Debian)

EXPERIENCE

Lead Engineer

QuestCraft (AI-first EdTech Platform)

August 2024 – Present

Kollam, Kerala, India

- Architected a **3-phase AI Exam Generation engine** that fulfills complex allocation constraints (topic × type × difficulty) while ensuring chunk-level diversity and 100% unique question variants.
- Engineered a **Finite State Automata (FSA)** approach for the Exam Review system, establishing a robust, state-based approval workflow for end-semester exams to ensure data integrity and process compliance.
- Developed a resilient material ingestion pipeline supporting chunked uploads (PDF/Video) with **pre-hash duplicate detection** and OCR, achieving an **85% reduction in memory consumption**.
- Led the migration from Huey to Celery + RabbitMQ, implementing multi-queue prioritization and memory-aware auto-restarts, resulting in **50% faster build speeds** and **20% lower infra overhead**.
- Implemented secure multi-tenant authentication using Logto with custom middleware for role-aware API context handling and token-based authorization.
- Integrated automated coding-question validation via Piston API and established performance monitoring using Langfuse and PostHog to track LLM accuracy and user engagement.
- Standardized infrastructure with Docker and Coolify on self-hosted VPS, managing automated periodic backups and optimizing Django ORM for high-concurrency production environments.
- **Tech Stack:** Django, Celery, RabbitMQ, Langchain, ChromaDB, Docker, Nginx, MinIO, PostgreSQL

LLM Researcher & Backend Developer

Traboda CyberLabs

November 2023 – August 2024

Hybrid – Kollam, India

- Led LLM research and prompt engineering for Prodago, boosting domain expert efficiency by 30% in generating AI governance playbooks.
- Developed scalable APIs using Azure Functions and FastAPI, resolving long-running task issues through Azure Durable Functions and batch processing optimization.
- Implemented LLMops pipelines with Agiflow and Promptfoo for model evaluation, reducing iteration cycles by 25%.
- **Skills:** FastAPI, Azure Functions, Pinecone, MLflow, OpenAI API, Claude API, LiteLLM

MLOps Engineer

TexNano

July 2022 – November 2023

Remote

- Designed MLOps infrastructure for a deep learning tool focused on pancreatic tumor diagnosis, streamlining model deployment and training workflows.
- Conducted research on novel CNN/Transformer approaches to enhance diagnostic accuracy for endoscopic ultrasound-guided aspirations.
- **Skills:** Python, Bash, Docker, CI/CD, Linux Server Administration

Associate Technology Engineer

Spoon

June 2021 – August 2021

Remote

- Developed serverless backend infrastructure using Firebase Cloud Functions and GCP, improving API response times by 20% through asynchronous invocations.
- Integrated KYC verification and billing systems, ensuring data security via Skyflow data privacy vaulting.

LEADERSHIP AND MENTORSHIP

inCTF | Core Organiser

May 2020 – June 2023

- Led a team of volunteers to execute one of India's largest hacking competitions, managing logistics and venue coordination for international participants.

amFOSS | Member and Mentor

August 2019 – Present

- Mentored junior developers in open-source contributions and full-stack development practices.